

# Predicting Forest Fires in Montesinho National Park – Portugal

*Gabriel Estivalet*

**DS-SF-24**

Aug/2016

# Getting the data

Coordinates in a 10x10 grid

Numerical ratings for the level of moisture in different ground types

Relative humidity (%)

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.0
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.0
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.0

Celsius

km/h

ml

Expected rate of fire spread

Burned area of a forest fire in ha

len(data) = 517

# Data Dictionary



**area:** Burned area of a forest fire (ha) 0-1190 ha

**X:** x-axis coordinate of the Montesinho park map: 1 to 9

**Y:** u-axis coordinate of the Montesinho park map: 2 to 9

**FFMC:** A numerical rating of the moisture content of litter and other cured fine fuels: 18.7 to 96.2

**DMC:** A numerical rating of the average moisture content of loosely compacted organic layers and medium-size woody material: 1.1 to 291.3

**DC:** A numerical rating of the average moisture content of deep, compact, organic layers: 7.9 to 860.6

**ISI:** A numerical rating of the expected rate of fire spread: 0.0 to 56.10

**Month:** month of the year: 1 to 12

**Day:** day of the week: 1 to 7

**Temp:** temperature in Celsius degrees: 2.2 to 33.30

**RH:** relative humidity in %: 15.0 to 100

**Wind:** wind speed in km/h: 0.40 to 9.40

**Rain:** outside rain in mm/m2: 0.0 to 6.4

How can we predict if there will be a fire?

If there is a fire, what's the size of the area that is most likely be affected?

What are the most important variables that can help with our predictions?

## Part I: Initial Analysis

# Initial Analysis

Created dummy variables for month and day

```
data = Dummify(data, ['month'], del_prms = True)  
data = Dummify(data, ['day'], del_prms = True)
```



	X	Y	FFMC	DMC	DC	ISI	temp	RH	wind	rain	..	month_may	month_nov	month_oct	month_sep	day_mon	day_sat	day_sun	day_thu	..
0	7	5	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	..	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	7	4	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	..	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0
2	7	4	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	..	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0


3 rows x 30 columns

# Initial Analysis

## Vizualizations - BoxPlots

```
plt.figure(figsize = (25,6))
plt.boxplot([data['FFMC'], data['DMC'], data['DC'], data['ISI'],
             data['temp'], data['RH'], data['wind'], data['rain']])

plt.xticks([1,2,3,4,5,6,7,8], ['FFMC', 'DMC', 'DC', 'ISI', 'temp', 'RH', 'wind', 'rain'], fontsize = 18)
plt.show()
```

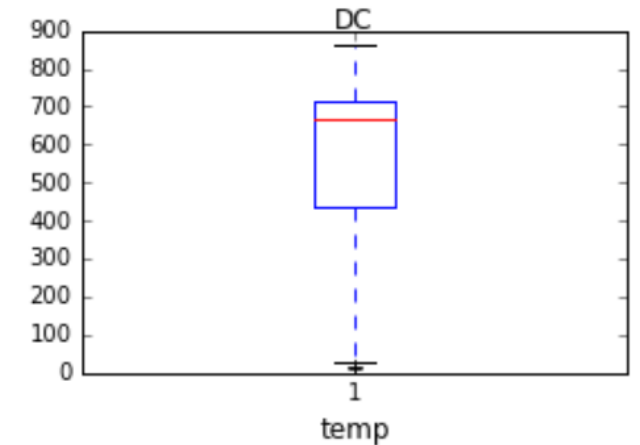
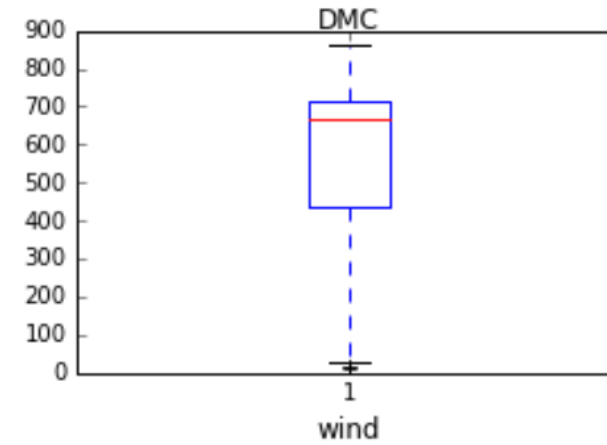








➡ The range of values was too big to display all box plots together!

# Initial Analysis

Separate box plots

?



ISI

RH

wind

temp



# Initial Analysis

```
data[ 'rain' ].value_counts() #Very dry place
```

```
0.0      509
```

```
0.8         2
```

```
0.2         2
```

```
0.4         1
```

```
6.4         1
```

```
1.4         1
```

```
1.0         1
```

```
Name: rain, dtype: int64
```

The values for rain are  
in milliliters.

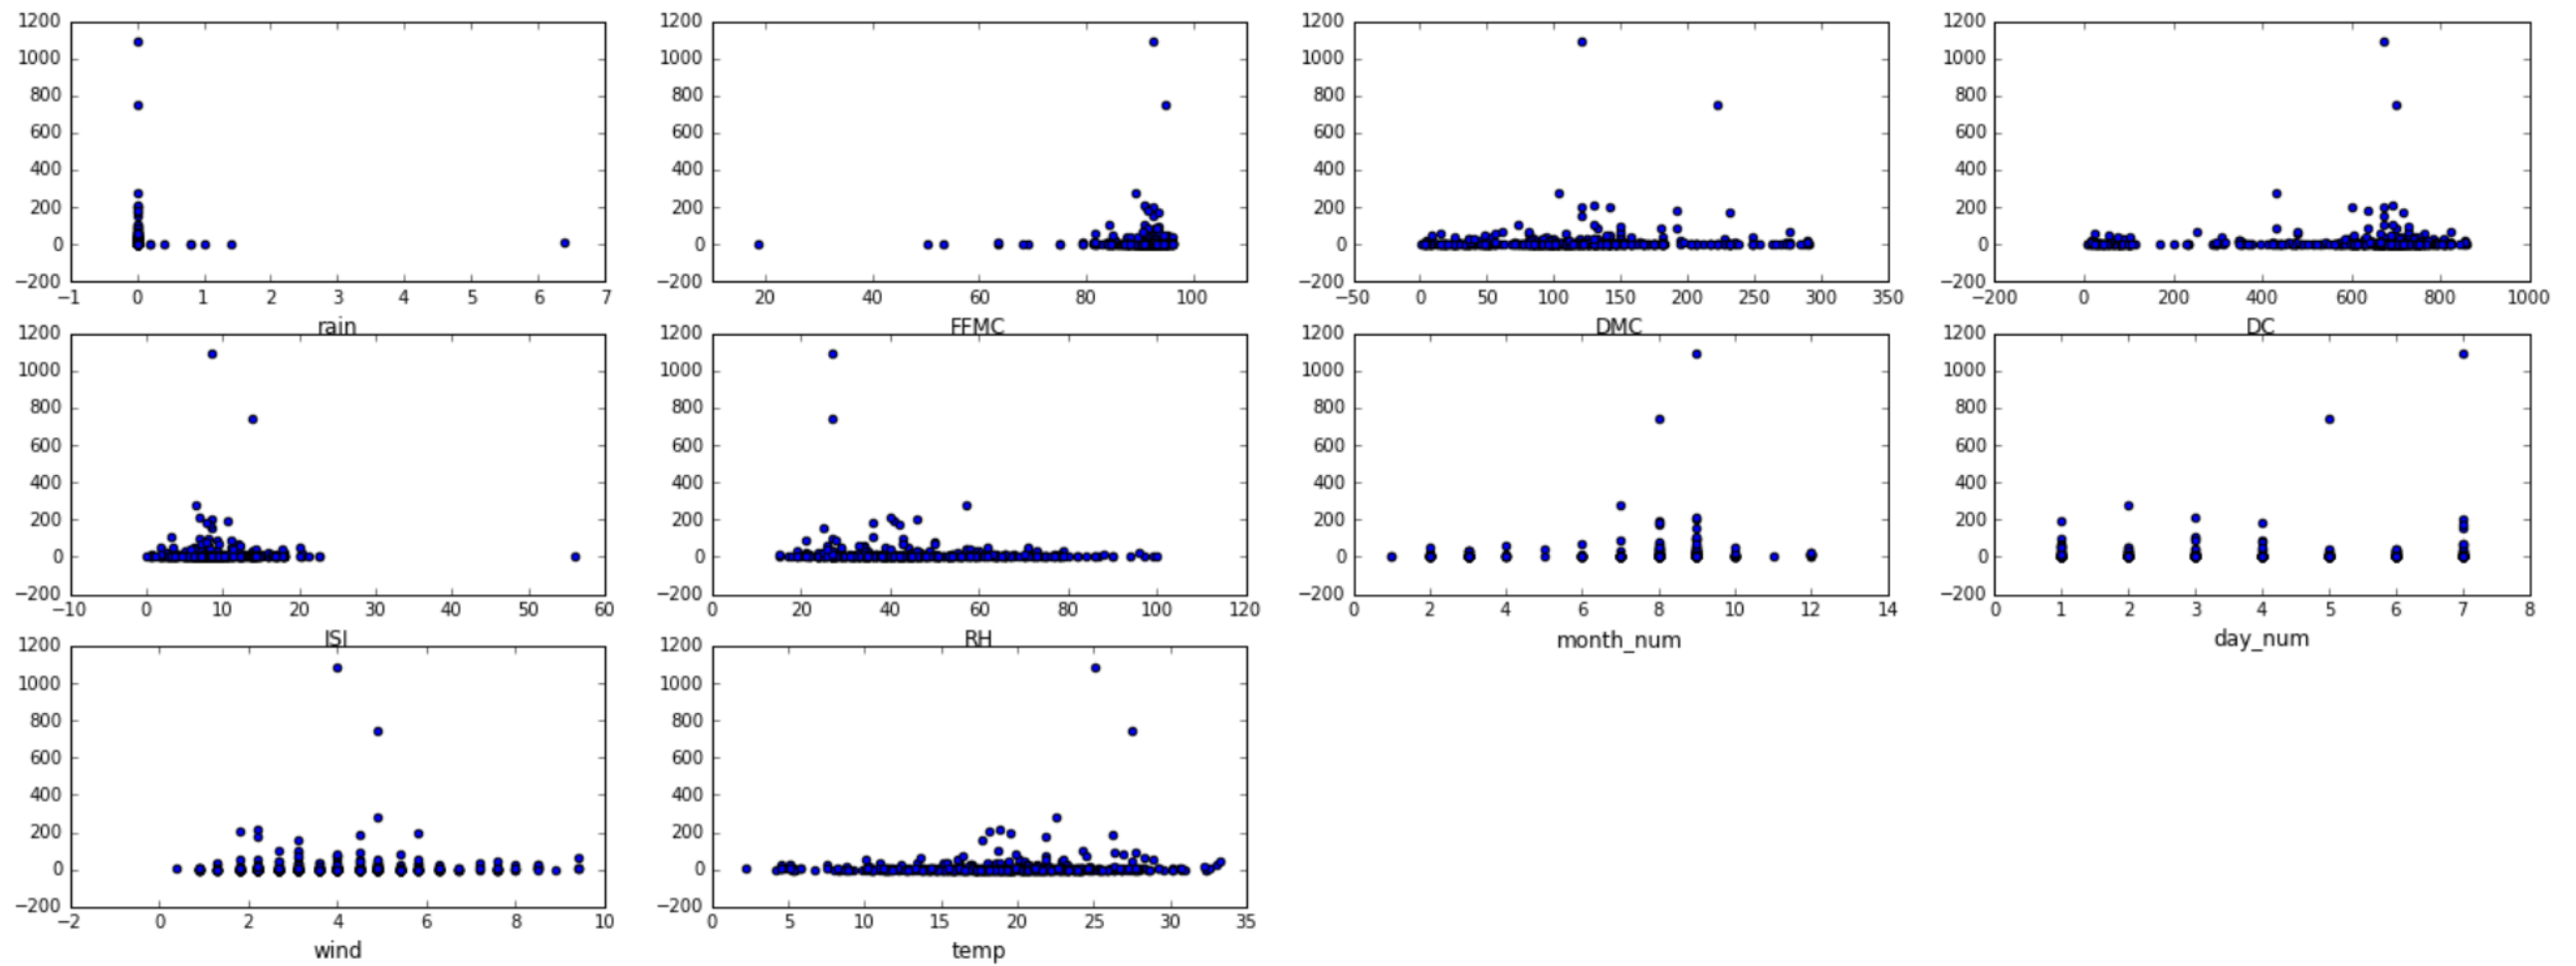


From 517 observations,  
509 had no rain at all.

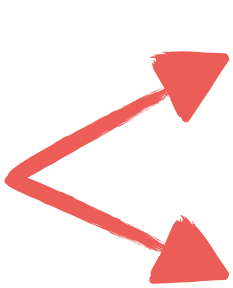
All models were used with and without "rain" and the  
results were approximately the same.

# Initial Analysis

## Scatter Plots



# Initial Analysis

Defined X and Y  y for Regression  
yclass for Classification

```
data["fire_yn"] = 0
data.fire_yn.loc[data['area'] > 0] = 1

ListOfAllVariables = data.columns.values
#print(ListOfAllVariables)
X = data[['FFMC', 'DMC', 'DC', 'ISI', 'temp', 'RH', 'wind', 'month_aug', 'month_dec', #9
          'month_feb', 'month_jan', 'month_jul', 'month_jun', 'month_mar', 'month_may', #6
          'month_nov', 'month_oct', 'month_sep', 'rain', 'day_mon', 'day_sat', 'day_sun', #7
          'day_thu', 'day_tue', 'day_wed']] #3
y = data['area']
yclass = data['fire_yn']
```

\*The variable "rain" did not influenced the results.

## Part II: Predicting a fire

# Classification

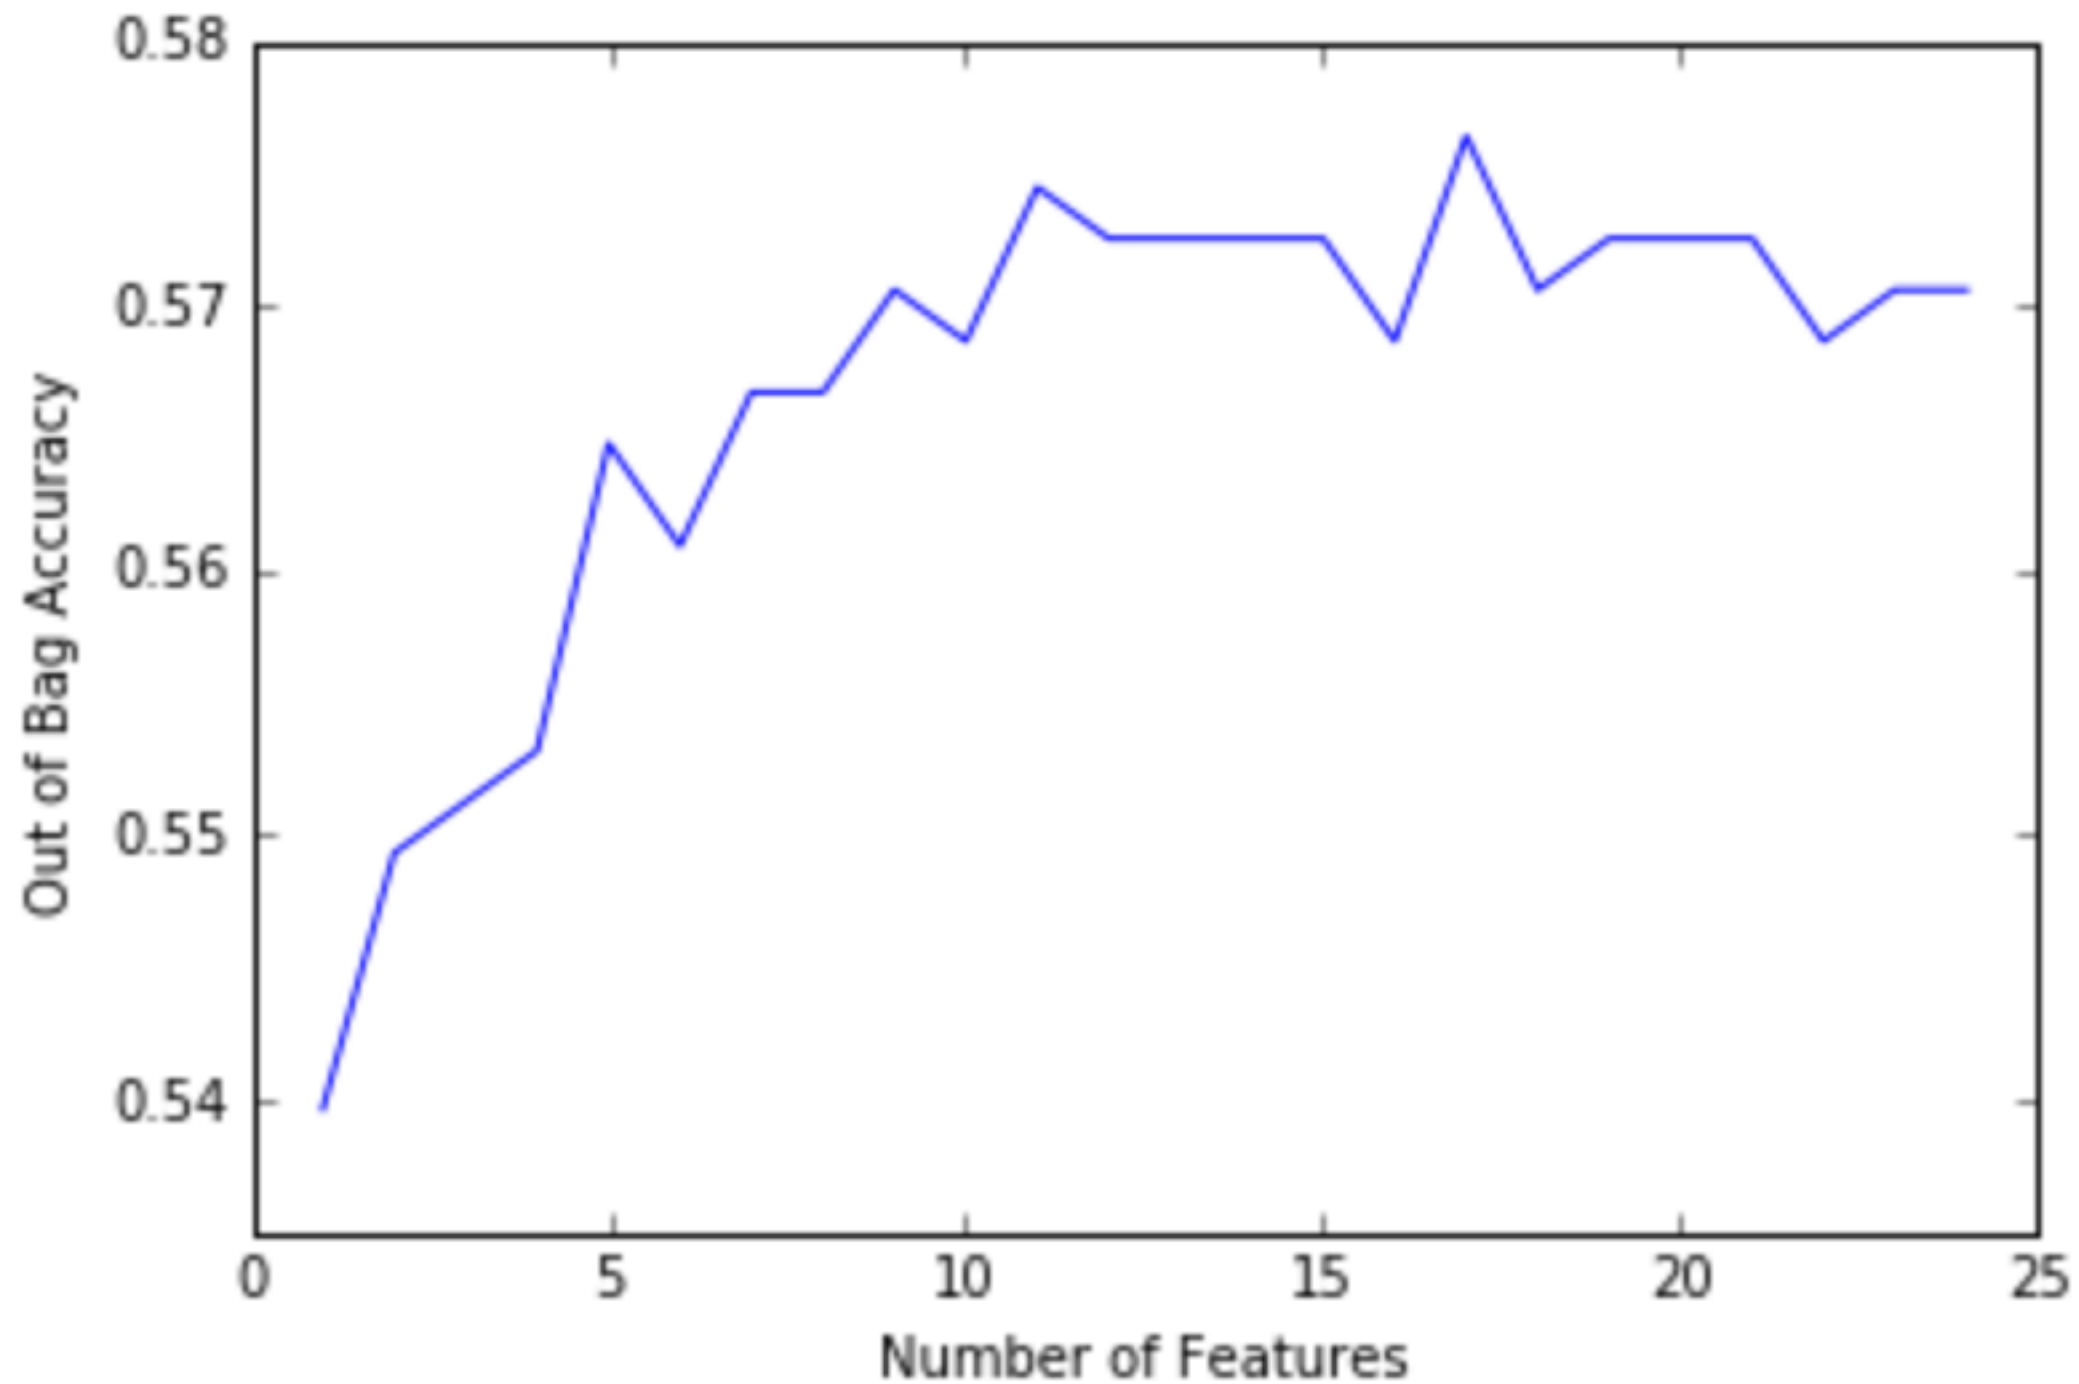
## Random Forest Classifier

```
Features = range(1,25)
oob_score_RF = []
for i in Features:
    RFCClass = RandomForestClassifier(n_estimators = 10000,
                                     max_features = i,
                                     min_samples_leaf = 5,
                                     oob_score = True,
                                     random_state = 1,
                                     n_jobs = -1)
    RFCClass.fit(X,yclass)
    oob_score_RF.append(RFCClass.oob_score_)

plt.plot(Features, oob_score_RF)
plt.xlabel("Number of Features")
plt.ylabel("Out of Bag Accuracy")
plt.show()

print("Out of Bag Accuracy = %f" %RFCClass.oob_score_)
scores = cross_val_score(RFCClass, X, yclass, cv = 10)
print("Cross-validation Accuracy = %f" %scores.mean())
```

# Classification




Out of Bag Accuracy = 0.570600

Cross-validation Accuracy = 0.448831

# Classification

## Multiple + Voting

```
#clf1 = LogisticRegression()  
clf2 = RandomForestClassifier(max_depth = 5, n_estimators = 10000)  
clf3 = BernoulliNB()  
clf4 = neighbors.KNeighborsClassifier(n_neighbors=199, weights='uniform')  
clf5 = GaussianNB()
```



\*Based on several attempts to run classification and regression models.

In both cases (voting = hard and soft),  
KNN seems to perform slightly better

# Classification

## Boosting Classification

```
NumberOfTrees = [100,1000,5000,10000,20000]
for i in NumberOfTrees:
    GBC_Tree = GradientBoostingClassifier(learning_rate = 0.01,
                                          n_estimators = i,
                                          max_depth = 2,
                                          min_samples_leaf = 10,
                                          random_state = 1)

kf = cross_validation.KFold(len(data), n_folds = 10, shuffle = True)

scores = []

for train_index, test_index in kf:
    GBC_Tree.fit(X.iloc[train_index], yclass.iloc[train_index])
    y_hat_test = GBC_Tree.predict(X.iloc[test_index])
    scores.append(float(sum(y_hat_test == yclass.iloc[test_index]))/len(y_hat_test))

Score_GBC_CV = np.mean(scores)

print(Score_GBC_CV)

0.560935143288
```

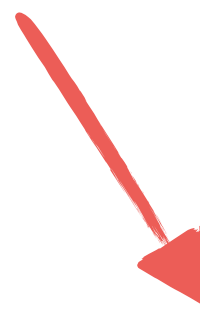


# Classification


## Confusion Matrix

```
#y_hat = RFClass.predict(X)
y_hat = GBC_Tree.predict(X)
confusion_matrix(yclass, y_hat)
```


```
array([[232, 15],
       [ 20, 250]])
```



False Positive Rate = 0.074074  
False Negative Rate = 0.060729  
True Positive Rate = 0.939271  
True Negative Rate = 0.925926  
Misclassification Error = 0.067698  
→ Accuracy = 0.932302



# Classification



## Part III: Predicting the affected area

# Regression

**Park area = 75000 ha**

Decision Trees → 65.62 ha

Random Forest → 64.32 ha


KNN → 64.53 ha

Boosting → 68.4 ha



68% of the time my prediction is within 68.4 ha from reality.  
95% of the time my prediction is within 136.8 ha from reality.

# Regression



```
(0.01980225268314029, 'month_aug'),  
(0.025151129716881708, 'month_sep'),  
(0.057454739710264456, 'day_sat'),  
(0.077926134041150327, 'wind'),  
(0.082412547758201449, 'ISI'),  
(0.089339478901834218, 'DC'),  
(0.1025578919855132, 'FFMC'),  
(0.12466673555985694, 'DMC'),  
(0.12607468830846508, 'RH'),  
(0.16540980955026477, 'temp')]
```

Could suggest that  
some fires are caused  
by humans in holidays

```
(0.0, 'rain'),
```

# Regression

## KNN

```
CV_Scores = []
RangeOfK = range(1,200)


for k in RangeOfK:
    knn = neighbors.KNeighborsRegressor(n_neighbors = k, weights = 'uniform')
    CV_Scores.append( -cross_val_score(knn, X, y, cv=10,
                                      scoring = 'mean_squared_error').mean())

plt.plot(RangeOfK, CV_Scores)
plt.xlabel("K")
plt.ylabel("Cross-Validation MSE")
plt.show()

print "The best K is %f" %RangeOfK[np.argmin(CV_Scores)]
```

# Regression

KNN



→ The best K is 199.000000

# Conclusion

---

- The best classification model can predict with an accuracy of 93.2% if there will be a fire.
- On average, the regression models can precise the area that will most likely be affected by a fire with an error of 64 ha 68% of the time and with an error of 128 ha 95% of the time.
- Temperature is the most important factor, followed by relative humidity and moisture levels. There may be cases when humans were behind the incidents.



Thank you!

# Sources

---

- ➔ Lichman, M. (2013). UCI Machine Learning Repository [<http://archive.ics.uci.edu/ml>]. Irvine, CA: University of California, School of Information and Computer Science.
- ➔ [Cortez and Morais, 2007] P. Cortez and A. Morais. A Data Mining Approach to Predict Forest Fires using Meteorological Data. In J. Neves, M. F. Santos and J. Machado Eds., New Trends in Artificial Intelligence, Proceedings of the 13th EPIA 2007 - Portuguese Conference on Artificial Intelligence, December, Guimarães, Portugal, pp. 512-523, 2007. APPIA, ISBN-13 978-989-95618-0-9. Available at: [\[Web Link\]](#)

# Appendix

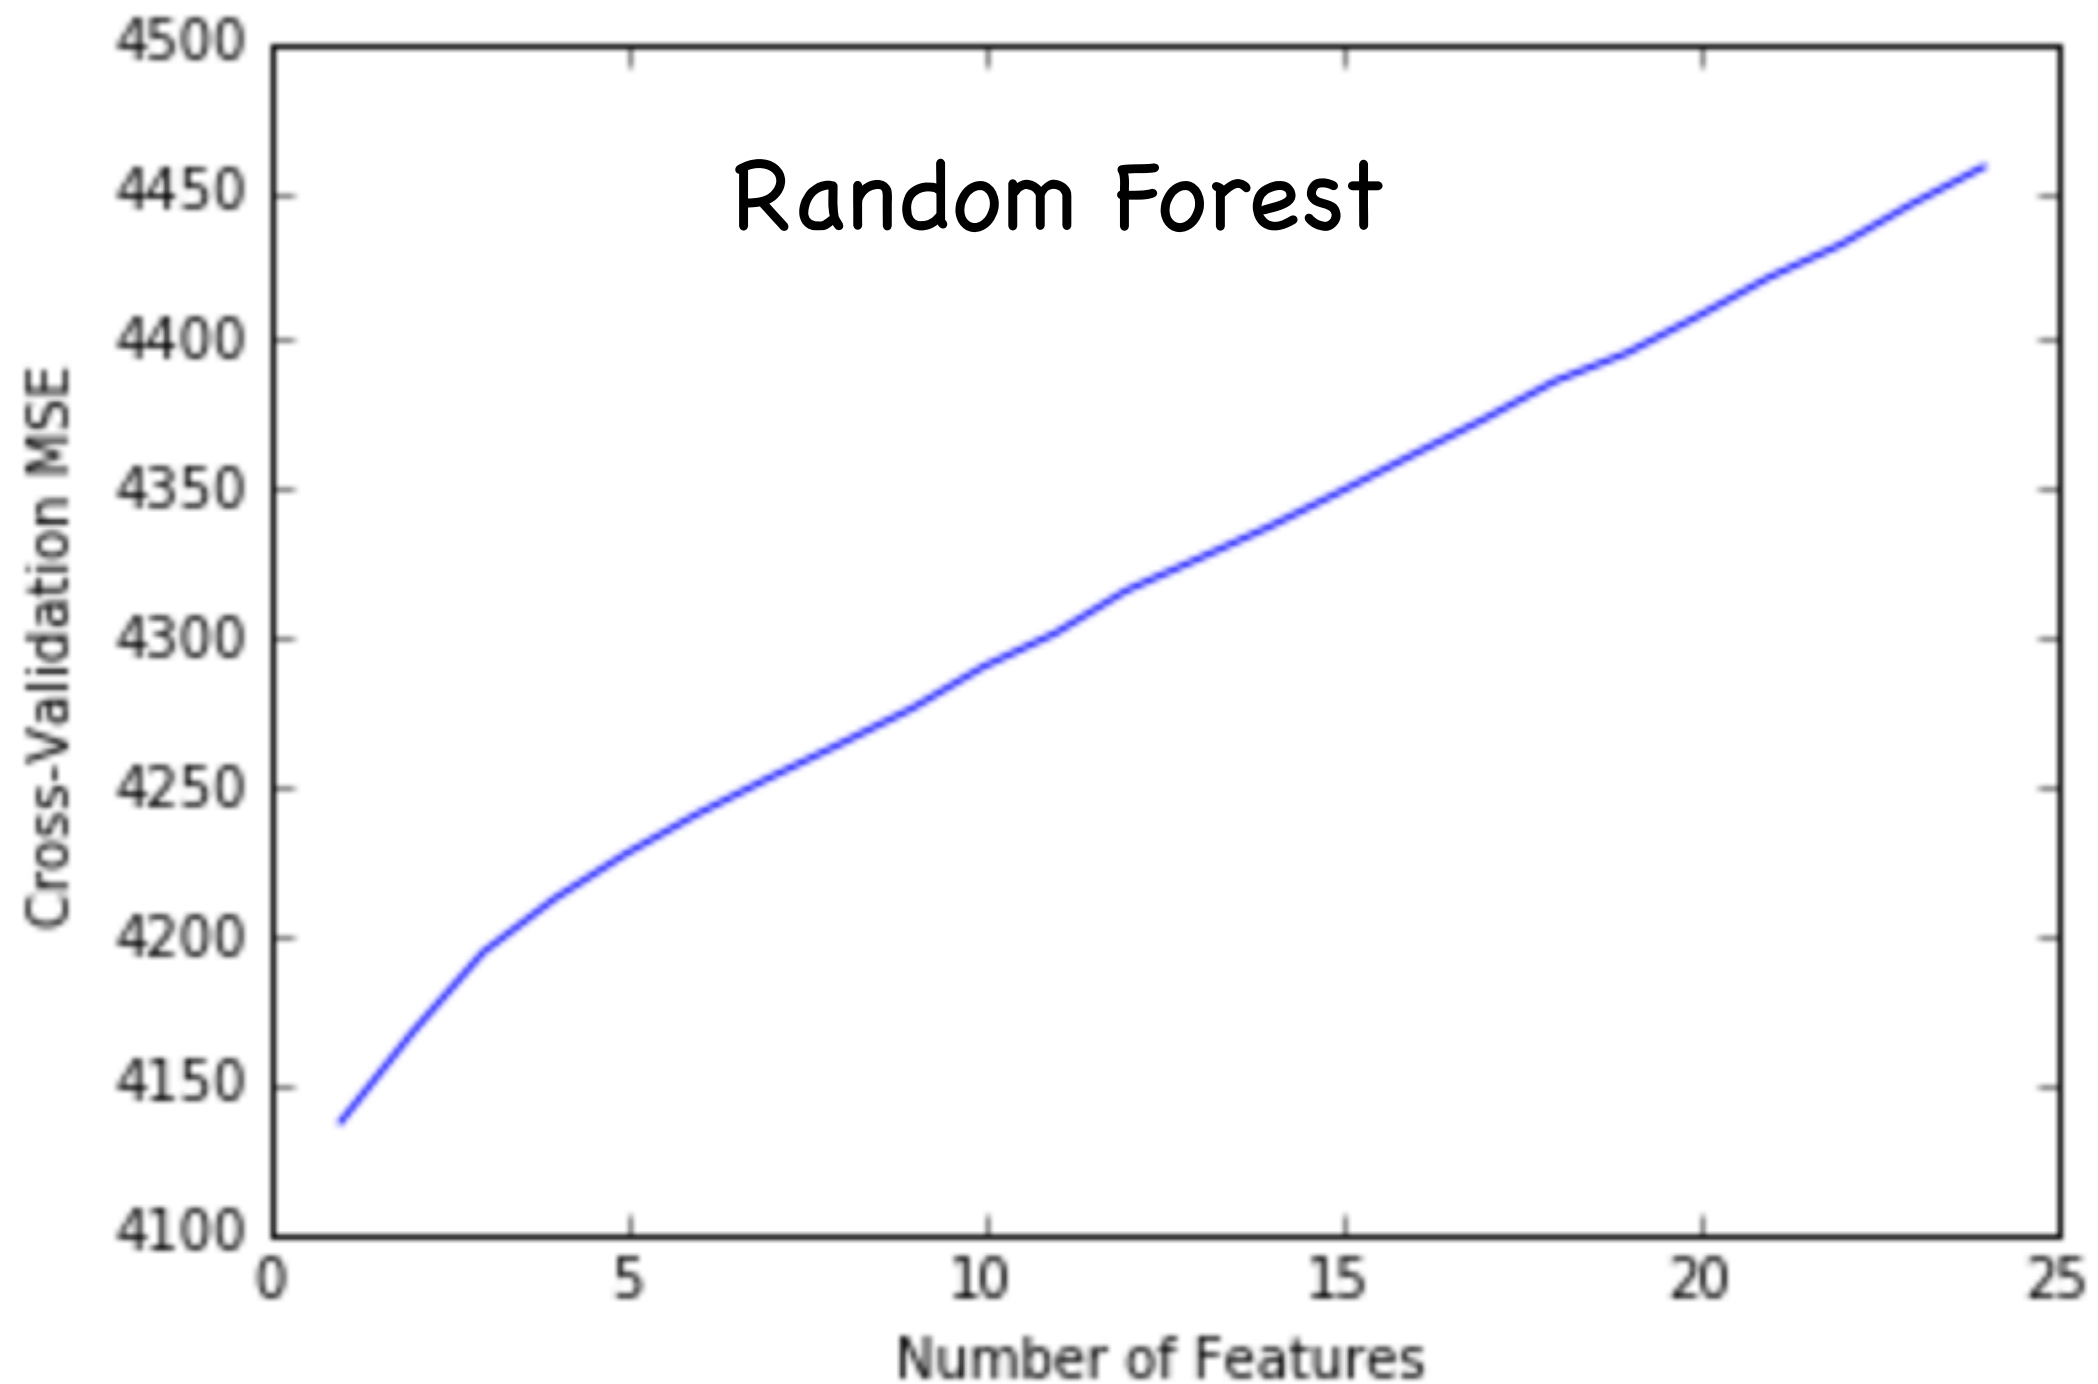
# Regression

## Random Forest

```
Features = range(1,25)
score = []
for i in Features:
    RF = RandomForestRegressor(n_estimators = 10000,
                              max_features = i,
                              min_samples_leaf = 10,
                              oob_score = True,
                              random_state = 1,
                              n_jobs = -1)
    score.append(-cross_val_score(RF, X, y, cv = 10,
                                  scoring = 'mean_squared_error',
                                  n_jobs = -1).mean())

plt.plot(Features, score)
plt.xlabel("Number of Features")
plt.ylabel("Cross-Validation MSE")
plt.show()
```

# Regression



Optimal level of features is 1, which leads to a MSE of 4137.37

# Regression

## Boosting

```
Score = []
NumberOfTrees = [100, 1000, 5000, 10000, 20000, 30000, 40000, 50000]
for i in NumberOfTrees:
    GBR_Tree = GradientBoostingRegressor(learning_rate = 0.01,
                                          n_estimators = i,
                                          max_depth = 2,
                                          min_samples_leaf = 10)

    Score.append(-cross_val_score(GBR_Tree, X, y, cv = 10,
                                   scoring = 'mean_squared_error', n_jobs = -1).mean()))

plt.plot(NumberOfTrees, Score)
plt.xlabel("Number of Trees")
plt.ylabel("CV - MSE")
plt.show()
```

# Regression

## Boosting

